

Task 3: Build a decision tree classifier to predict whether a customer will purchase a product or service based on their demographic and behavioral data. Use a dataset such as the Bank Marketing dataset from the UCI Machine Learning Repository.

Import Libraries

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import matplotlib.pyplot as plt
```

Load Dataset

```
❶ data = pd.read_csv("healthcare_dataset.csv")
data.head(5)
```

---	Name	Age	Gender	Blood Type	Medical Condition	Date of Admission	Doctor	Hospital	Insurance Provider	Billing Amount	Room Number	Admission Type	Discharge Date	Medication	Test Results
0	Bobby JacksOn	30	Male	B-	Cancer	1/31/2024	Matthew Smith	Sons and Miller	Blue Cross	18856.28131	328	Urgent	2/2/2024	Paracetamol	Normal
1	LesLieTErRy	62	Male	A+	Obesity	8/20/2019	Samantha Davies	Kim Inc	Medicare	33643.32729	265	Emergency	8/28/2019	Ibuprofen	Inconclusive
2	DaNnYsMith	78	Female	A-	Obesity	9/22/2022	Tiffany Mitchell	Cook PLC	Aetna	27955.09808	205	Emergency	10/7/2022	Aspirin	Normal
3	andrew waTiS	28	Female	O+	Diabetes	11/18/2020	Kevin Wells	Hernandez Rogers and Vang,	Medicare	37909.78241	450	Elective	12/18/2020	Ibuprofen	Abnormal
4	adriENNE bEll	43	Female	AB+	Cancer	9/19/2022	Kathleen Hanna	White-White	Aetna	14238.31781	458	Urgent	10/9/2022	Penicillin	Abnormal

Basic Data Check

```
❷ data.info()
```

```
*** <class 'pandas.core.frame.DataFrame'>
RangeIndex: 55500 entries, 0 to 55499
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Name            55500 non-null   object  
 1   Age             55500 non-null   int64  
 2   Gender          55500 non-null   object  
 3   Blood Type     55500 non-null   object  
 4   Medical Condition 55500 non-null   object  
 5   Date of Admission 55500 non-null   object  
 6   Doctor          55500 non-null   object  
 7   Hospital         55500 non-null   object  
 8   Insurance Provider 55500 non-null   object  
 9   Billing Amount   55500 non-null   float64 
 10  Room Number     55500 non-null   int64  
 11  Admission Type  55500 non-null   object  
 12  Discharge Date  55500 non-null   object  
 13  Medication       55500 non-null   object  
 14  Test Results    55500 non-null   object  
dtypes: float64(1), int64(2), object(12)
memory usage: 6.4+ MB
```

```
data.isnull().sum()
```

	0
Name	0
Age	0
Gender	0
Blood Type	0
Medical Condition	0
Date of Admission	0
Doctor	0
Hospital	0
Insurance Provider	0
Billing Amount	0
Room Number	0
Admission Type	0
Discharge Date	0
Medication	0
Test Results	0

dtype: int64

Data Cleaning

Fill Missing Values

```
for col in data.columns:
    if data[col].dtype == 'object':
        data[col].fillna(data[col].mode()[0], inplace=True)
    else:
        data[col].fillna(data[col].mean(), inplace=True)

... /tmp/ipython-input-834479767.py:3: FutureWarning: A value is trying to
The behavior will change in pandas 3.0. This inplace method will never
For example, when doing 'df[col].method(value, inplace=True)', try usi

        data[col].fillna(data[col].mode()[0], inplace=True)
/tmp/ipython-input-834479767.py:5: FutureWarning: A value is trying to
The behavior will change in pandas 3.0. This inplace method will never
For example, when doing 'df[col].method(value, inplace=True)', try usi

        data[col].fillna(data[col].mean(), inplace=True)
```

Remove Duplicate Rows

```
data.drop_duplicates(inplace=True)
```

Select Features and Target

```
# Admission Type  
X = data[['Age', 'Gender', 'Blood Type', 'Medical Condition',  
          'Insurance Provider', 'Billing Amount']]  
  
y = data['Admission Type']
```

Convert Categorical Columns into Numbers

```
X = pd.get_dummies(X)  
y = pd.get_dummies(y)
```

Train-Test Split

```
▶ X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=42  
)
```

Train Decision Tree Model

```
model = DecisionTreeClassifier(max_depth=5)  
model.fit(X_train, y_train)  
  
+ DecisionTreeClassifier ⓘ ⓘ  
DecisionTreeClassifier(max_depth=5)
```

Prediction

```
y_pred = model.predict(X_test)
```

Model Evaluation

Accuracy

```
] accuracy = accuracy_score(y_test, y_pred)  
print("Model Accuracy:", accuracy)  
  
Model Accuracy: 0.00227396761870111
```

Confusion Matrix

```
▶ confusion_matrix(y_test.values.argmax(axis=1),  
    y_pred.argmax(axis=1))  
  
... array([[3657,      5,      3],  
          [3677,      5,      9],  
          [3626,      7,      5]])
```

Classification Report

```
▶ print(classification_report(  
    y_test.values.argmax(axis=1),  
    y_pred.argmax(axis=1)  
)  
  
...      precision    recall   f1-score   support  
  
       0       0.33     1.00     0.50     3665  
       1       0.29     0.00     0.00     3691  
       2       0.29     0.00     0.00     3638  
  
accuracy                           0.33    10994  
macro avg       0.31     0.33     0.17    10994  
weighted avg     0.31     0.33     0.17    10994
```

Visualize Decision Tree

```
▶ plt.figure(figsize=(20,10))  
plot_tree(model,  
          feature_names=X.columns,  
          filled=True)  
plt.show()  
  
...  
  
...
```

